

# VAE

Variational autoencoder [1] models inherit autoencoder architecture, but use variational approach  
homework, we will implement VAE and quantitatively measure the quality of the generated sample

[1] Auto-Encoding Variational Bayes, Diederik P Kingma, Max Welling 2013 <https://arxiv.org/abs/1312.6114>

[2] Improved techniques for training gans, Salimans, T., Goodfellow, I., Zaremba, W., Cheung, V., Radford, A. 2016  
Neural Information Processing Systems

[3] A note on inception score, Shane Barratt, Rishi Sharma 2018 <https://arxiv.org/abs/1801.01973>

## ▼ PART I. Train a good VAE model

### ▼ Setup

```
import tensorflow as tf
if tf.__version__ < '2.0.0':
    tf.enable_eager_execution()
tf.executing_eagerly()

import numpy as np
import os

import matplotlib.pyplot as plt
import matplotlib.gridspec as gridspec

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# A bunch of utility functions

def show_images(images):
    # images reshape to (batch_size, D)
    images = np.reshape(images, [images.shape[0], -1])
    sqrt_n = int(np.ceil(np.sqrt(images.shape[0])))
    sqrt_m = int(np.ceil(np.sqrt(images.shape[1])))

    fig = plt.figure(figsize=(sqrt_n, sqrt_m))
    gs = gridspec.GridSpec(sqrt_n, sqrt_m)
    gs.update(wspace=0.05, hspace=0.05)

    for i, img in enumerate(images):
        ax = plt.subplot(gs[i])
        plt.axis('off')
        ax.set_xticklabels([])
```

```

        ax.set_yticklabels([])
        ax.set_aspect('equal')
        plt.imshow(img.reshape([sqrt(img), sqrt(img)]))
    return

def preprocess_img(x):
    return 2 * x - 1.0

def rel_error(x,y):
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))

def count_params(model):
    """Count the number of parameters in the current TensorFlow graph """
    param_count = np.sum([np.prod(p.shape) for p in model.weights])
    return param_count

```

## ▼ Dataset

We will be working on the MNIST dataset, which is 60,000 training and 10,000 test images. Each p digit on black background (0 through 9). This was one of the first datasets used to train convolutic standard CNN model can easily exceed 99% accuracy.

**Heads-up:** Our MNIST wrapper returns images as vectors. That is, they're size (batch, 784). If you v resize them to (batch,28,28) or (batch,28,28,1). They are also type np.float32 and bounded [0,1].

```

class MNIST(object):
    def __init__(self, batch_size, shuffle=False):
        """
        Construct an iterator object over the MNIST data

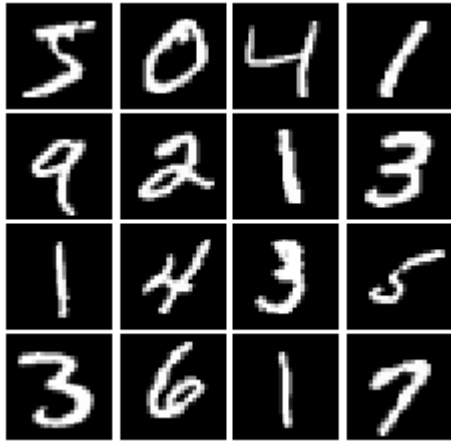
        Inputs:
        - batch_size: Integer giving number of elements per minibatch
        - shuffle: (optional) Boolean, whether to shuffle the data on each epoch
        """
        train, _ = tf.keras.datasets.mnist.load_data()
        X, y = train
        X = X.astype(np.float32)/255
        X = X.reshape((X.shape[0], -1))
        self.X, self.y = X, y
        self.batch_size, self.shuffle = batch_size, shuffle

    def __iter__(self):
        N, B = self.X.shape[0], self.batch_size
        idxs = np.arange(N)
        if self.shuffle:
            np.random.shuffle(idxs)
        return iter((self.X[i:i+B], self.y[i:i+B]) for i in range(0, N, B))

# show a batch
mnist = MNIST(batch_size=16)
show_images(mnist.X[:16])

```

📄 Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>  
 11493376/11490434 [=====] - 0s 0us/step



```
X_DIM = mnist.X[0].size
num_samples = 100000
num_to_show = 100

# Hyperparamters. Your job to find these.
# TODO:
# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
num_epochs = 100
batch_size = 100
Z_DIM = 5
learning_rate = 5e-4
# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
```

## ▼ Encoder

Our first step is to build a variational encoder network  $q_\phi(z | x)$ .

**Hint:** You should use the layers in `tf.keras.layers` to build the model. Use four FC layers. All fully c  
 For initialization, just use the default initializer used by the `tf.keras.layers` functions.

The output of the encoder should thus have shape `[batch_size, 2*z_dim]`, and contain real number diagonal log variance  $\log \sigma(x_i)^2$  of each of the `batch_size` input images. Note, we want to make i stability.

**WARNING:** Do not apply any non-linearity to the last activation.

```
def q_phi(z_dim=Z_DIM, x_dim=X_DIM):
    model = tf.keras.models.Sequential([
        # TODO: implement architecture
        # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
        tf.keras.layers.Dense(392, activation="relu", use_bias=True, input_shape=(x_dim,)),
        tf.keras.layers.Dense(196, activation="relu", use_bias=True),
        tf.keras.layers.Dense(128, activation="tanh", use_bias=True),
        tf.keras.layers.Dense(2 * z_dim, use_bias=True)
        # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
    ])
    return model
```

```
# TODO: implement reparameterization trick
def sample_z(mu, log_var):
    # Your code here for the reparameterization trick.
    # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
    samples = None
    #print(mu.shape)
    z = tf.random.normal(tf.shape(mu))
    s = tf.math.exp(0.5 * log_var)
    samples = mu + s * z
    return samples
    # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
```

## ▼ Decoder

Now to build a decoder network  $p_\theta(x | z)$ . You should use the layers in `tf.keras.layers` to construct. Connected layers should include bias terms. Note that you can use the `tf.nn` module to access activation functions and initializers for parameters.

In this exercise, we will use Bernoulli MLP decoder where  $p_\theta(x | z)$  is modeled with multivariate Bernoulli distribution we discussed in the lecture, as following (see Appendix C.1 in the original paper).

$$\log p(x | z) = \sum_{i=1} x_i \log z_i + (1 - x_i) \log(1 - z_i)$$

Note, the output of the decoder should have shape `[batch_size, x_dim]` and should output the unnormalized probabilities.

**WARNING:** Do not apply any non-linearity to the last activation.

```
def p_theta(z_dim=Z_DIM, x_dim=X_DIM):
    model = tf.keras.models.Sequential([
        # TODO: implement architecture
        # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
        tf.keras.layers.Dense(128, activation="tanh", use_bias=True, input_shape=(z_dim,)),
        tf.keras.layers.Dense(196, activation="relu", use_bias=True),
        tf.keras.layers.Dense(392, activation="relu", use_bias=True),
        tf.keras.layers.Dense(x_dim, use_bias=True)

        # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
    ])
    return model
```

## ▼ Loss definition

Compute the VAE loss.

1. For the reconstruction loss, you might find `tf.nn.sigmoid_cross_entropy_with_logits` or `tf.keras.losses.binary_crossentropy`.
2. For the kl loss, we discussed the closed form kl divergence between two gaussians in the lecture.

```
def vae_loss(x, x_logit, z_mu, z_logvar):
    recon_loss = None
    kl_loss = None # negative value
```

```

# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
#print(x, x_logit)
#bce = tf.keras.losses.BinaryCrossentropy(from_logits=True, reduction=tf.keras.losses.Reduction.N
recon_loss = tf.reduce_sum(tf.nn.sigmoid_cross_entropy_with_logits(x, x_logit), axis=1)
#entropy = bce(x, x_logit)
#recon_loss = tf.reduce_sum(entropy)
#print(entropy.shape, recon_loss.shape)
temp = 1 + z_logvar - tf.square(z_mu) - tf.math.exp(z_logvar)
kl_loss = -0.5 * (tf.reduce_sum(temp, axis = 1))
#print(kl_loss, z_mu, z_logvar)
#print(recon_loss, kl_loss)
# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
#print(kl_loss)
vae_loss = tf.reduce_mean(recon_loss + kl_loss)
return vae_loss, tf.reduce_mean(recon_loss)

```

## ▼ Optimizing our loss

```

Q = q_phi()
P = p_theta()
solver = tf.keras.optimizers.Adam(learning_rate)
mnist = MNIST(batch_size=batch_size, shuffle=True)

```


### Visualize generated samples before training

```

z_gen = tf.random.normal(shape=[num_to_show, Z_DIM])
x_gen = P(z_gen)
imgs_numpy = tf.nn.sigmoid(x_gen).numpy()
show_images(imgs_numpy)
plt.show()

```





## ▼ Training a VAE!

If everything works, your batch average reconstruction loss should drop below 95.



```

iter_count = 0
show_every = 400
for epoch in range(num_epochs):
    for (x_i, _) in mnist:
        with tf.GradientTape() as tape:
            z_concat = Q(preprocess_img(x_i))
            z_mu, z_logvar = tf.split(z_concat, num_or_size_splits=2, axis=1)
            z_i = sample_z(z_mu, z_logvar)

            x_logit = P(z_i)
            loss, recon_loss = vae_loss(x_i, x_logit, z_mu, z_logvar)

        grads = tape.gradient(loss,
                               [Q.trainable_variables, P.trainable_variables])

        solver.apply_gradients(zip([*grads[0], *grads[1]],
                                   [*Q.trainable_variables, *P.trainable_variables]))

    if (iter_count % show_every == 0):
        print('Epoch: {}, Iter: {}, Loss: {:.4}, Recon: {:.4}'.format(
            epoch, iter_count, loss, recon_loss))
        #imgs_numpy = tf.nn.sigmoid(x_logit).numpy()
        #show_images(imgs_numpy[0:16])
        #plt.show()
    iter_count += 1

```



Epoch: 0, Iter: 0, Loss: 546.7, Recon: 544.4  
Epoch: 0, Iter: 400, Loss: 143.0, Recon: 133.9  
Epoch: 1, Iter: 800, Loss: 146.4, Recon: 136.4  
Epoch: 2, Iter: 1200, Loss: 129.1, Recon: 118.9  
Epoch: 2, Iter: 1600, Loss: 127.4, Recon: 117.2  
Epoch: 3, Iter: 2000, Loss: 136.3, Recon: 125.4  
Epoch: 4, Iter: 2400, Loss: 123.0, Recon: 112.2  
Epoch: 4, Iter: 2800, Loss: 122.5, Recon: 111.7  
Epoch: 5, Iter: 3200, Loss: 131.5, Recon: 120.2  
Epoch: 6, Iter: 3600, Loss: 120.8, Recon: 109.6  
Epoch: 6, Iter: 4000, Loss: 120.3, Recon: 109.4  
Epoch: 7, Iter: 4400, Loss: 127.7, Recon: 116.2  
Epoch: 8, Iter: 4800, Loss: 117.1, Recon: 105.7  
Epoch: 8, Iter: 5200, Loss: 118.6, Recon: 107.3  
Epoch: 9, Iter: 5600, Loss: 124.6, Recon: 113.0  
Epoch: 10, Iter: 6000, Loss: 115.2, Recon: 103.4  
Epoch: 10, Iter: 6400, Loss: 116.0, Recon: 104.6  
Epoch: 11, Iter: 6800, Loss: 123.6, Recon: 111.9  
Epoch: 12, Iter: 7200, Loss: 114.6, Recon: 102.7  
Epoch: 12, Iter: 7600, Loss: 114.9, Recon: 103.3  
Epoch: 13, Iter: 8000, Loss: 120.5, Recon: 108.5  
Epoch: 14, Iter: 8400, Loss: 113.0, Recon: 101.0  
Epoch: 14, Iter: 8800, Loss: 114.3, Recon: 102.6  
Epoch: 15, Iter: 9200, Loss: 119.7, Recon: 107.7  
Epoch: 16, Iter: 9600, Loss: 112.2, Recon: 100.2  
Epoch: 16, Iter: 10000, Loss: 113.0, Recon: 101.1  
Epoch: 17, Iter: 10400, Loss: 119.1, Recon: 107.1  
Epoch: 18, Iter: 10800, Loss: 110.1, Recon: 98.11  
Epoch: 18, Iter: 11200, Loss: 112.6, Recon: 100.7  
Epoch: 19, Iter: 11600, Loss: 117.2, Recon: 105.0  
Epoch: 20, Iter: 12000, Loss: 110.3, Recon: 98.27  
Epoch: 20, Iter: 12400, Loss: 111.5, Recon: 99.46  
Epoch: 21, Iter: 12800, Loss: 116.9, Recon: 104.6  
Epoch: 22, Iter: 13200, Loss: 110.1, Recon: 97.97  
Epoch: 22, Iter: 13600, Loss: 111.6, Recon: 99.47  
Epoch: 23, Iter: 14000, Loss: 116.0, Recon: 103.7  
Epoch: 24, Iter: 14400, Loss: 109.2, Recon: 96.86  
Epoch: 24, Iter: 14800, Loss: 110.4, Recon: 98.39  
Epoch: 25, Iter: 15200, Loss: 115.6, Recon: 103.2  
Epoch: 26, Iter: 15600, Loss: 109.0, Recon: 96.77  
Epoch: 26, Iter: 16000, Loss: 109.9, Recon: 97.56  
Epoch: 27, Iter: 16400, Loss: 114.8, Recon: 102.4  
Epoch: 28, Iter: 16800, Loss: 108.7, Recon: 96.25  
Epoch: 28, Iter: 17200, Loss: 109.8, Recon: 97.72  
Epoch: 29, Iter: 17600, Loss: 114.0, Recon: 101.6  
Epoch: 30, Iter: 18000, Loss: 109.1, Recon: 96.74  
Epoch: 30, Iter: 18400, Loss: 109.2, Recon: 96.76  
Epoch: 31, Iter: 18800, Loss: 113.6, Recon: 101.0  
Epoch: 32, Iter: 19200, Loss: 108.4, Recon: 95.97  
Epoch: 32, Iter: 19600, Loss: 109.1, Recon: 96.91  
Epoch: 33, Iter: 20000, Loss: 112.9, Recon: 100.2  
Epoch: 34, Iter: 20400, Loss: 107.8, Recon: 95.33  
Epoch: 34, Iter: 20800, Loss: 109.0, Recon: 96.74  
Epoch: 35, Iter: 21200, Loss: 112.0, Recon: 99.29  
Epoch: 36, Iter: 21600, Loss: 107.4, Recon: 94.9  
Epoch: 36, Iter: 22000, Loss: 108.5, Recon: 96.07  
Epoch: 37, Iter: 22400, Loss: 113.3, Recon: 100.6  
Epoch: 38, Iter: 22800, Loss: 107.2, Recon: 94.72

```
Epoch: 38, Iter: 23200, Loss: 109.2, Recon: 96.74
Epoch: 39, Iter: 23600, Loss: 113.0, Recon: 100.3
Epoch: 40, Iter: 24000, Loss: 107.7, Recon: 95.02
Epoch: 40, Iter: 24400, Loss: 108.0, Recon: 95.67
Epoch: 41, Iter: 24800, Loss: 112.1, Recon: 99.38
Epoch: 42, Iter: 25200, Loss: 106.8, Recon: 94.33
Epoch: 42, Iter: 25600, Loss: 107.8, Recon: 95.49
Epoch: 43, Iter: 26000, Loss: 112.9, Recon: 100.1
Epoch: 44, Iter: 26400, Loss: 106.8, Recon: 94.21
Epoch: 44, Iter: 26800, Loss: 108.4, Recon: 95.91
Epoch: 45, Iter: 27200, Loss: 111.7, Recon: 98.76
Epoch: 46, Iter: 27600, Loss: 106.6, Recon: 94.01
Epoch: 46, Iter: 28000, Loss: 107.6, Recon: 95.2
Epoch: 47, Iter: 28400, Loss: 112.3, Recon: 99.28
Epoch: 48, Iter: 28800, Loss: 106.7, Recon: 94.18
Epoch: 48, Iter: 29200, Loss: 107.7, Recon: 95.3
Epoch: 49, Iter: 29600, Loss: 111.2, Recon: 98.14
Epoch: 50, Iter: 30000, Loss: 107.7, Recon: 95.09
Epoch: 50, Iter: 30400, Loss: 106.7, Recon: 94.34
Epoch: 51, Iter: 30800, Loss: 111.4, Recon: 98.47
Epoch: 52, Iter: 31200, Loss: 106.4, Recon: 93.87
Epoch: 52, Iter: 31600, Loss: 107.5, Recon: 95.04
Epoch: 53, Iter: 32000, Loss: 111.2, Recon: 98.25
Epoch: 54, Iter: 32400, Loss: 105.9, Recon: 93.23
Epoch: 54, Iter: 32800, Loss: 107.0, Recon: 94.57
Epoch: 55, Iter: 33200, Loss: 111.4, Recon: 98.49
Epoch: 56, Iter: 33600, Loss: 106.5, Recon: 93.89
Epoch: 56, Iter: 34000, Loss: 106.9, Recon: 94.37
Epoch: 57, Iter: 34400, Loss: 110.3, Recon: 97.25
Epoch: 58, Iter: 34800, Loss: 106.1, Recon: 93.51
Epoch: 58, Iter: 35200, Loss: 107.3, Recon: 94.81
Epoch: 59, Iter: 35600, Loss: 110.0, Recon: 97.03
Epoch: 60, Iter: 36000, Loss: 106.2, Recon: 93.64
Epoch: 60, Iter: 36400, Loss: 106.5, Recon: 93.87
Epoch: 61, Iter: 36800, Loss: 111.4, Recon: 98.36
Epoch: 62, Iter: 37200, Loss: 105.5, Recon: 92.96
Epoch: 62, Iter: 37600, Loss: 107.1, Recon: 94.58
Epoch: 63, Iter: 38000, Loss: 110.2, Recon: 97.3
Epoch: 64, Iter: 38400, Loss: 106.8, Recon: 94.11
Epoch: 64, Iter: 38800, Loss: 106.6, Recon: 94.08
Epoch: 65, Iter: 39200, Loss: 109.6, Recon: 96.68
Epoch: 66, Iter: 39600, Loss: 104.9, Recon: 92.21
Epoch: 66, Iter: 40000, Loss: 106.2, Recon: 93.76
Epoch: 67, Iter: 40400, Loss: 109.1, Recon: 96.09
Epoch: 68, Iter: 40800, Loss: 105.8, Recon: 93.19
Epoch: 68, Iter: 41200, Loss: 106.8, Recon: 94.1
Epoch: 69, Iter: 41600, Loss: 109.7, Recon: 96.83
Epoch: 70, Iter: 42000, Loss: 106.1, Recon: 93.46
Epoch: 70, Iter: 42400, Loss: 106.4, Recon: 93.77
Epoch: 71, Iter: 42800, Loss: 109.4, Recon: 96.42
Epoch: 72, Iter: 43200, Loss: 105.7, Recon: 92.98
Epoch: 72, Iter: 43600, Loss: 106.7, Recon: 94.1
Epoch: 73, Iter: 44000, Loss: 109.7, Recon: 96.59
Epoch: 74, Iter: 44400, Loss: 104.5, Recon: 91.67
Epoch: 74, Iter: 44800, Loss: 106.6, Recon: 93.86
Epoch: 75, Iter: 45200, Loss: 110.4, Recon: 97.36
Epoch: 76, Iter: 45600, Loss: 104.4, Recon: 91.66
Epoch: 76, Iter: 46000, Loss: 106.0, Recon: 93.37
```



```
Epoch: 77, Iter: 46400, Loss: 109.3, Recon: 96.33
Epoch: 78, Iter: 46800, Loss: 104.6, Recon: 91.86
Epoch: 78, Iter: 47200, Loss: 106.4, Recon: 93.68
Epoch: 79, Iter: 47600, Loss: 109.3, Recon: 96.31
Epoch: 80, Iter: 48000, Loss: 104.8, Recon: 92.17
Epoch: 80, Iter: 48400, Loss: 105.9, Recon: 93.37
Epoch: 81, Iter: 48800, Loss: 109.6, Recon: 96.44
Epoch: 82, Iter: 49200, Loss: 104.6, Recon: 91.97
Epoch: 82, Iter: 49600, Loss: 106.1, Recon: 93.5
Epoch: 83, Iter: 50000, Loss: 108.5, Recon: 95.31
Epoch: 84, Iter: 50400, Loss: 104.5, Recon: 91.88
Epoch: 84, Iter: 50800, Loss: 105.5, Recon: 92.79
Epoch: 85, Iter: 51200, Loss: 108.1, Recon: 94.98
Epoch: 86, Iter: 51600, Loss: 105.0, Recon: 92.19
Epoch: 86, Iter: 52000, Loss: 105.7, Recon: 93.06
Epoch: 87, Iter: 52400, Loss: 109.0, Recon: 95.72
Epoch: 88, Iter: 52800, Loss: 105.3, Recon: 92.65
Epoch: 88, Iter: 53200, Loss: 105.5, Recon: 92.92
```

### Visualize generated samples after training

```
Epoch: 90, Iter: 54400, Loss: 105.6, Recon: 93.04
z_gen = tf.random.normal(shape=[num_to_show, Z_DIM])
x_gen = P(z_gen)
imgs_numpy = tf.nn.sigmoid(x_gen).numpy()
show_images(imgs_numpy)
plt.show()
```





## ▼ PART II. Compute the inception score for your trained VAE m

In this part, we will quantitatively measure how good your VAE model is.



### ▼ Train a classifier

We first need to train a classifier.



```
batch_size = 128
num_classes = 10
epochs = 20

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()

x_train = x_train.reshape(60000, 784)
x_test = x_test.reshape(10000, 784)
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
y_train = tf.keras.utils.to_categorical(y_train, num_classes)
y_test = tf.keras.utils.to_categorical(y_test, num_classes)

model = tf.keras.models.Sequential()
model.add(tf.keras.layers.Dense(512, activation='relu', input_shape=(784,)))
model.add(tf.keras.layers.Dropout(0.2))
model.add(tf.keras.layers.Dense(512, activation='relu'))
model.add(tf.keras.layers.Dropout(0.2))
model.add(tf.keras.layers.Dense(num_classes, activation='softmax'))

model.summary()

model.compile(loss='categorical_crossentropy',
              optimizer=tf.keras.optimizers.RMSprop(),
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

60000 train samples  
 10000 test samples  
 Model: "sequential\_16"

Layer (type)	Output Shape	Param #
dense_64 (Dense)	(None, 512)	401920
dropout (Dropout)	(None, 512)	0
dense_65 (Dense)	(None, 512)	262656
dropout_1 (Dropout)	(None, 512)	0
dense_66 (Dense)	(None, 10)	5130

Total params: 669,706  
 Trainable params: 669,706  
 Non-trainable params: 0

```
Epoch 1/20
469/469 [=====] - 2s 4ms/step - loss: 0.2448 - accuracy: 0.9259 - va
Epoch 2/20
469/469 [=====] - 2s 4ms/step - loss: 0.1008 - accuracy: 0.9695 - va
Epoch 3/20
469/469 [=====] - 2s 4ms/step - loss: 0.0748 - accuracy: 0.9774 - va
Epoch 4/20
469/469 [=====] - 2s 4ms/step - loss: 0.0595 - accuracy: 0.9814 - va
Epoch 5/20
469/469 [=====] - 2s 4ms/step - loss: 0.0513 - accuracy: 0.9848 - va
Epoch 6/20
469/469 [=====] - 2s 3ms/step - loss: 0.0442 - accuracy: 0.9866 - va
Epoch 7/20
469/469 [=====] - 2s 4ms/step - loss: 0.0375 - accuracy: 0.9888 - va
Epoch 8/20
469/469 [=====] - 2s 4ms/step - loss: 0.0341 - accuracy: 0.9900 - va
Epoch 9/20
469/469 [=====] - 2s 3ms/step - loss: 0.0298 - accuracy: 0.9914 - va
Epoch 10/20
469/469 [=====] - 2s 4ms/step - loss: 0.0278 - accuracy: 0.9915 - va
Epoch 11/20
469/469 [=====] - 2s 3ms/step - loss: 0.0277 - accuracy: 0.9923 - va
Epoch 12/20
469/469 [=====] - 2s 3ms/step - loss: 0.0242 - accuracy: 0.9928 - va
Epoch 13/20
469/469 [=====] - 2s 3ms/step - loss: 0.0246 - accuracy: 0.9931 - va
Epoch 14/20
469/469 [=====] - 2s 4ms/step - loss: 0.0217 - accuracy: 0.9939 - va
Epoch 15/20
469/469 [=====] - 2s 4ms/step - loss: 0.0219 - accuracy: 0.9937 - va
Epoch 16/20
469/469 [=====] - 2s 4ms/step - loss: 0.0183 - accuracy: 0.9945 - va
Epoch 17/20
469/469 [=====] - 2s 4ms/step - loss: 0.0193 - accuracy: 0.9946 - va
Epoch 18/20
469/469 [=====] - 2s 4ms/step - loss: 0.0165 - accuracy: 0.9953 - va
Epoch 19/20
469/469 [=====] - 2s 4ms/step - loss: 0.0189 - accuracy: 0.9949 - va
```

Epoch 20/20


469/469 [=====] - 2s 3ms/step - loss: 0.0179 - accuracy: 0.9958 - va

Test loss: 0.14113247394561768

## ▼ Verify the trained classifier on the generated samples

Generate samples and visually inspect if the predicted labels on the samples match the actual dig

```
z_gen = tf.random.normal(shape=[num_samples, Z_DIM])
x_gen = P(z_gen)
imgs_numpy = tf.nn.sigmoid(x_gen[:num_to_show]).numpy()
show_images(imgs_numpy)
plt.show()
```



```
np.argmax(model.predict(tf.nn.sigmoid(x_gen[:20])), axis=-1)
```



```
array([9, 6, 8, 8, 7, 0, 4, 5, 1, 8, 8, 8, 1, 3, 1, 4, 0, 8, 8, 9])
```

## ▼ Implement the inception score

Implement Equation 1 in the reference [3]. Replace expectation in the equation with empirical average and exponentiation at the end. You should get Inception score of at least 9.0.

```

kld_obj = tf.keras.losses.KLDivergence()
# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

score = None
image = tf.nn.sigmoid(x_gen)
predicted_y = model.predict(image) #p(y|x)
true_y = np.ones((num_samples, 1)) * np.mean(predicted_y, axis = 0)

score = np.exp(kld_obj(predicted_y, true_y))

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
print('Inception score: {:.4}'.format(score))

```

➞ Inception score: 9.229

## ▼ Plot the histogram of predicted labels

Let's additionally inspect the class diversity of the generated samples.

```

plt.hist(np.argmax(model.predict(tf.nn.sigmoid(x_gen)), axis=-1),
         bins=np.arange(11)-0.5, rwidth=0.8, density=True)
plt.xticks(range(10))
plt.show()

```

