

forward-backward-kl

June 7, 2020

0.1 KL divergence minimization

In this exercise, we will minimize KL divergence between the true distribution $p(x)$ and our model $q_\theta(x)$ using gradient descent. We will be utilizing [Tensorflow Probability \(TFP\)](#) built on Tensorflow that makes it easy to combine probabilistic models and deep learning. If you haven't already, please go ahead and [install](#) the library.

```
[0]: import numpy as np
import tensorflow.compat.v1 as tf
tf.disable_eager_execution()
import tensorflow_probability as tfp
import matplotlib.pyplot as plt
%matplotlib inline
import subprocess

FIGSIZE = (6, 6)

[0]: # Evaluate the probabilities given the samples
def prob_map(p, lim, N):
    x = np.linspace(lim[0], lim[1], N)
    y = np.linspace(lim[2], lim[3], N)
    xy = np.array(np.meshgrid(x, y)).reshape(2, N**2).T
    xy_prob = p.prob(xy.astype(np.float32)).eval()
    x, y = np.meshgrid(x, y)
    return x, y, xy_prob.reshape((N, N))

# Create mean and covariance params for bivariate normal distribution
def gen_gaussian_params(d=2):
    p_mean = np.random.standard_normal((d,)).astype(np.float32)
    p_cov = np.random.standard_normal((d, d)).astype(np.float32)
    p_cov[0, 1] = 0
    p_cov = np.dot(p_cov, p_cov.T)
    return p_mean, p_cov

# Plot helpers
def contour(x, y, xy_prob, lim, cmap=plt.cm.inferno, axis=None):
    if axis is None:
        plt.contour(x, y, xy_prob, np.linspace(0., 1.1*xy_prob.max(), 10),
                    zorder=2, extent=lim, linewidths=4, colors='w', alpha=0.5)
```

```

plt.contour(x, y, xy_prob, np.linspace(0., 1.1*xy_prob.max(), 10),
            zorder=3, extent=lim, linewidths=2, cmap=cmap)
else:
    axis.contour(x, y, xy_prob, np.linspace(0., 1.1*xy_prob.max(), 10),
                 zorder=2, extent=lim, linewidths=4, colors='w', alpha=0.5)
    axis.contour(x, y, xy_prob, np.linspace(0., 1.1*xy_prob.max(), 10),
                 zorder=3, extent=lim, linewidths=2, cmap=cmap)

def scatter(sample, axis=None):
    if axis is None:
        plt.scatter(sample[:200, 0], sample[:200, 1], s=25, alpha=0.7, c='k')
    else:
        axis.scatter(sample[:200, 0], sample[:200, 1], s=25, alpha=0.7, c='k')

```

0.1.1 Create a bivariate normal as the true distribution $p(x) = \mathcal{N}(\mu, \Sigma)$

Draw the probability contour plot for the true distribution $p(x)$

```

[3]: np.random.seed(555)
      tf.reset_default_graph()
      sess = tf.InteractiveSession()

      p_mean, p_cov = gen_gaussian_params()
      p_mean += 0.8
      p = tfp.distributions.MultivariateNormalFullCovariance(p_mean, p_cov)

      plt.figure(figsize=FIGSIZE)
      N = 128
      lim = [-2.5, 3.5, -2.5, 3.5]
      x, y, xy_prob = prob_map(p, lim, N)
      contour(x, y, xy_prob, lim, cmap=plt.cm.viridis)

      plt.axis(lim)
      plt.axis('off')
      plt.show()


```

```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-
packages/tensorflow_probability/python/distributions/distribution.py:332:
MultivariateNormalFullCovariance.__init__ (from
tensorflow_probability.python.distributions.mvn_full_covariance) is deprecated
and will be removed after 2019-12-01.
Instructions for updating:
`MultivariateNormalFullCovariance` is deprecated, use
`MultivariateNormalTriL(loc=loc,
scale_tril=tf.linalg.cholesky(covariance_matrix))` instead.
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-
packages/tensorflow/python/ops/linalg/linear_operator_lower_triangular.py:158:
calling LinearOperator.__init__ (from

```


tensorflow.python.ops.linalg.linear_operator) with graph_parents is deprecated and will be removed in a future version.
Instructions for updating:
Do not pass `graph_parents`. They will no longer be used.



Draw samples from the true distribution $p(x)$ and visualize

```
[0]: n_sample = 1000
      sample = p.sample(n_sample).eval()

      plt.figure(figsize=FIGSIZE)
      scatter(sample)
      plt.axis(lim)
      plt.axis('off')
      plt.show()
```



0.2 Part A. Forward KL minimization

0.2.1 Create the model distribution

Now we will create $q_\theta(x)$ and apply gradient descent to minimize the forward KL divergence

$$\theta^* = \arg \min_{\theta} D_{KL}(p(x) \parallel q_\theta(x))$$

```
[0]: # Define your model here!
qm = np.random.standard_normal((2,)).astype(np.float32)
qr = np.random.standard_normal((2, 2)).astype(np.float32)

q_mean = tf.Variable(qm) # will converge to true p
q_raw = tf.Variable(qr)
q_cov = tf.matmul(q_raw, tf.transpose(q_raw))

q = tfp.distributions.MultivariateNormalFullCovariance(q_mean, q_cov)
print(type(p_mean), type(q_mean))
# Define the loss here!
```

```

mu12 = tf.reshape(p_mean - q_mean, [1,2])
q_cov_inv = tf.linalg.inv(q_cov)
sum1 = tf.matmul(tf.matmul(mu12, q_cov_inv), tf.transpose(mu12))
trc = tf.linalg.trace(tf.matmul(q_cov_inv, p.covariance()))

loss = 0.5 * (tf.math.log(tf.linalg.det(q_cov)) + trc + sum1)
train = tf.train.GradientDescentOptimizer(0.1).minimize(loss, var_list = [
    →[q_mean, q_raw])
sess.run(tf.global_variables_initializer())

```

```

<class 'numpy.ndarray'> <class
'tensorflow.python.ops.resource_variable_ops.ResourceVariable'>

```

Lets run it for 20 iterations and visualize $q_{\theta(t)}(x)$

```

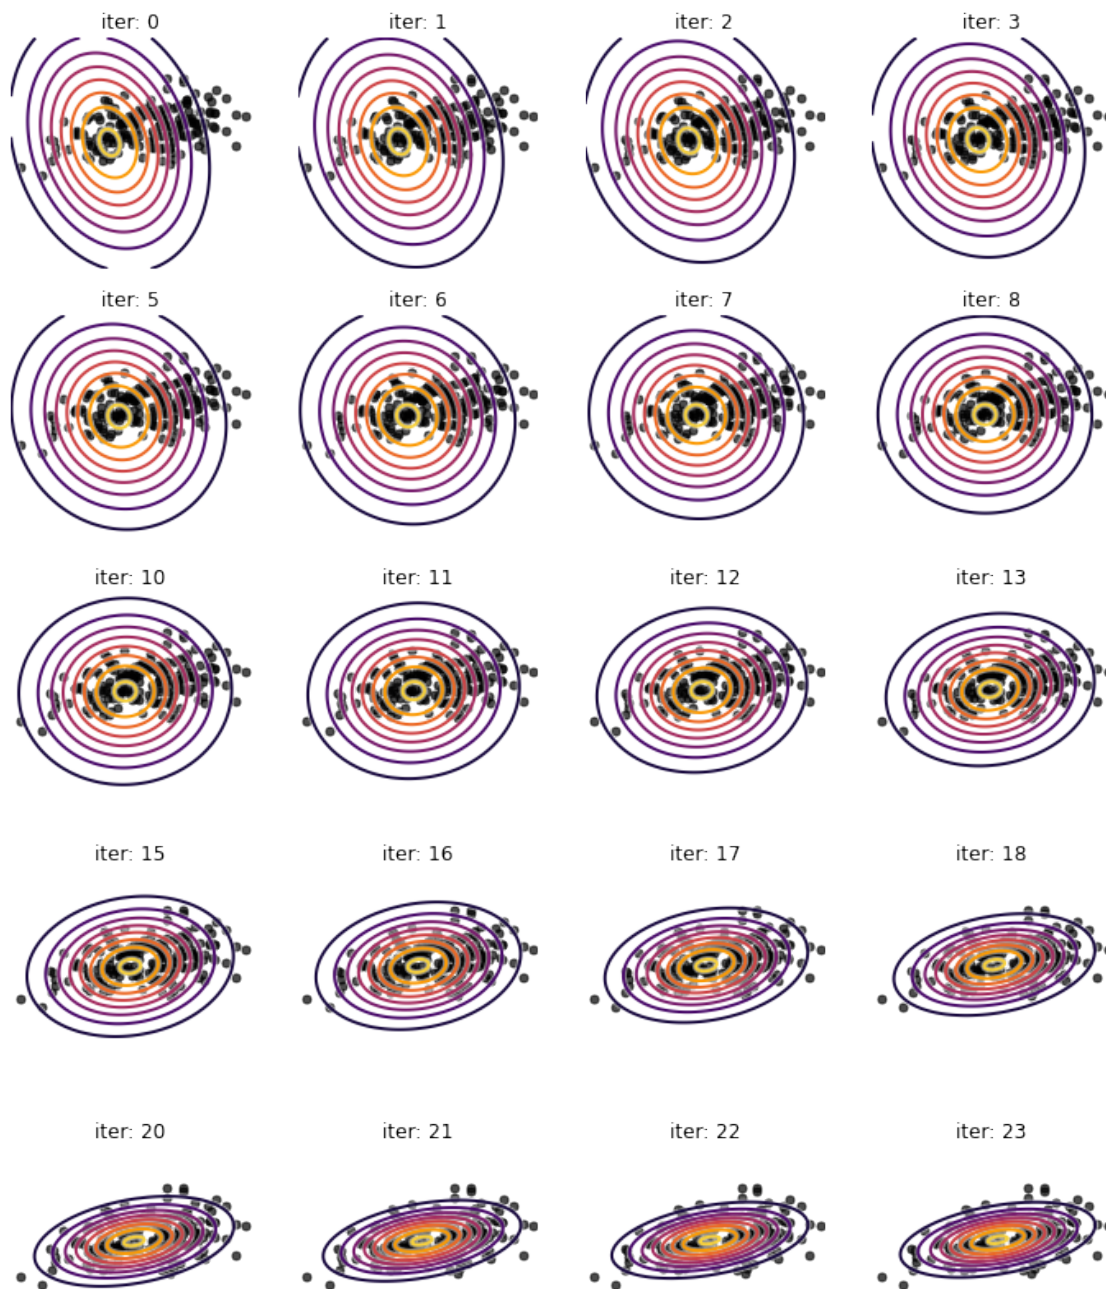
[0]: fig, axes = plt.subplots(5,4, figsize=[12, 15])

for i in range(5):
    for j in range(4):
        #print(q_raw.eval(session=sess), q_mean.eval(session=sess), q_temp.
        →eval(session=sess))
        scatter(sample, axis=axes[i][j])
        x, y, xy_prob = prob_map(q, lim, N)
        contour(x, y, xy_prob, lim, axis=axes[i][j])
        axes[i][j].axis(lim)
        axes[i][j].axis('off')
        axes[i][j].set_title('iter: {}'.format(i*5 + j))
        sess.run(train)

plt.show()
plt.close(fig)

sess.close()

```



0.2.2 Create a mixture of two bivariate normals as the true distribution

$p(x) = \pi_1 \mathcal{N}(\mu_1, \Sigma_1) + \pi_2 \mathcal{N}(\mu_2, \Sigma_2)$ where $\pi_1 + \pi_2 = 1$

```
[11]: def build_p(mix=0.8):
      return tfp.distributions.Mixture(
```

```

cat=tfp.distributions.Categorical(probs=[mix, 1.-mix]),
components=[
    tfp.distributions.MultivariateNormalFullCovariance(
        loc=[0.5, -0.7], covariance_matrix=[[1., 1.], [1., 2.]]),
    tfp.distributions.MultivariateNormalFullCovariance(
        loc=[-2.5, 2.5], covariance_matrix=[[0.5, 0.1], [0.1, 0.5]]),
])

np.random.seed(8)
tf.reset_default_graph()
sess = tf.InteractiveSession()

p = build_p()

```

```

/usr/local/lib/python3.6/dist-packages/tensorflow/python/client/session.py:1751:
UserWarning: An interactive session is already active. This can cause out-of-
memory errors in some cases. You must explicitly call
`InteractiveSession.close()` to release resources held by the other session(s).
  warnings.warn('An interactive session is already active. This can '


```

Draw the probability contour plot for the true distribution $p(x)$

```

[12]: N = 128
lim = [-4, 4, -4, 4]
plt.figure(figsize=FIGSIZE)
x, y, xy_prob = prob_map(p, lim, N)
contour(x, y, xy_prob, lim, cmap=plt.cm.viridis)
plt.axis(lim)
plt.axis('off')
plt.show()

```



Draw samples from the true distribution $p(x)$ and visualize

```
[13]: n_sample = 1000
      sample = p.sample(n_sample).eval()
      plt.figure(figsize=FIGSIZE)
      scatter(sample)
      plt.axis(lim)
      plt.axis('off')
      plt.show()
```




0.2.3 Create the model distribution

Now we will create $q_\theta(x)$ and apply gradient descent to minimize the forward KL divergence

$$\theta^* = \arg \min_{\theta} D_{KL}(p(x) \parallel q_{\theta}(x))$$

```
[0]: # Define your model here!
qm = np.random.standard_normal((2,)).astype(np.float32)
qr = np.random.standard_normal((2, 2)).astype(np.float32)

q_mean = tf.Variable(qm) # will converge to true p
q_raw = tf.Variable(qr)
q_cov = tf.matmul(q_raw, tf.transpose(q_raw))
#print(a.dtype,b,q_mean)
q = tfp.distributions.MultivariateNormalFullCovariance(q_mean, q_cov)

#print(sample.shape, p_mean.shape, p_cov.shape)

# Define the loss here!
```

```

loss = -tfp.monte_carlo.expectation(f = (lambda x : q.log_prob(x)), samples =
    ↳sample)
#logp = None
train = tf.train.GradientDescentOptimizer(0.1).minimize(loss, var_list =
    ↳[q_mean, q_raw])
sess.run(tf.global_variables_initializer())

```

Lets run it for 20 iterations and visualize $q_{\theta^{(t)}}(x)$

```


[15]: fig, axes = plt.subplots(5,4, figsize=[12, 15])

for i in range(5):
    for j in range(4):
        scatter(sample, axis=axes[i][j])
        x, y, xy_prob = prob_map(q, lim, N)
        contour(x, y, xy_prob, lim, axis=axes[i][j])
        axes[i][j].axis(lim)
        axes[i][j].axis('off')
        axes[i][j].set_title('iter: {}'.format(i*5 + j))
        sess.run(train)

plt.show()
plt.close(fig)

sess.close()

```



0.3 Part B. Reverse KL minimization

```
[16]: tf.reset_default_graph()
      sess = tf.InteractiveSession()
      p = build_p()
```

```
/usr/local/lib/python3.6/dist-packages/tensorflow/python/client/session.py:1751:
UserWarning: An interactive session is already active. This can cause out-of-
memory errors in some cases. You must explicitly call
`InteractiveSession.close()` to release resources held by the other session(s).
warnings.warn('An interactive session is already active. This can '
```

0.3.1 Create the model distribution

Now we will create $q_{\theta}(x)$ and apply gradient descent to minimize the reverse KL divergence

$$\theta^* = \arg \min_{\theta} D_{KL}(q_{\theta}(x) \parallel p(x))$$

```
[0]: # Define your model here!
qm = np.random.standard_normal((2,)).astype(np.float32)
qr = np.random.standard_normal((2, 2)).astype(np.float32)

q_mean = tf.Variable(qm) # will converge to true p
q_raw = tf.Variable(qr)
q_cov = tf.matmul(q_raw, tf.transpose(q_raw))

## 2 -> p, 1 -> q

q = tfp.distributions.MultivariateNormalFullCovariance(q_mean, q_cov)
loss = tfp.monte_carlo.expectation(
    f=lambda x: -p.log_prob(x) + q.log_prob(x),
    samples = q.sample(1000)
)

# Define the loss here!
train = tf.train.GradientDescentOptimizer(0.1).minimize(loss, var_list = [
    q_mean, q_raw])
sess.run(tf.global_variables_initializer())
```

Lets run it for 20 iterations and visualize $q_{\theta^{(t)}}(x)$

```
[19]: fig, axes = plt.subplots(5,4, figsize=[12, 15])

for i in range(5):
    for j in range(4):
        scatter(sample, axis=axes[i][j])
        x, y, xy_prob = prob_map(q, lim, N)
        contour(x, y, xy_prob, lim, axis=axes[i][j])
        axes[i][j].axis(lim)
        axes[i][j].axis('off')
        axes[i][j].set_title('iter: {}'.format(i*5 + j))
        sess.run(train)

plt.show()
plt.close(fig)
```

```
sess.close()
```

